


12 SEPTEMBER 2024



**Mdigital**

# Local LLMs & RAG

with Ollama and Postgres

---

# Local large language models

1. Why use local LLMs?
2. How to use Ollama
3. Local Vector DB
4. Put it together: RAG
5. If we have time: Fine-tuning

# Apologies

I'm a web developer...

---



**Windows  
and  
Python**



**macOS  
and  
JavaScript**

WHY?

# Local LLM advantages

- ✓ You can keep your data private
- ✓ You don't need to use paid APIs
- ✓ Efficient - limit power consumption
- ✓ Learning and experimentation

OLLAMA

# Ollama

Easy to install

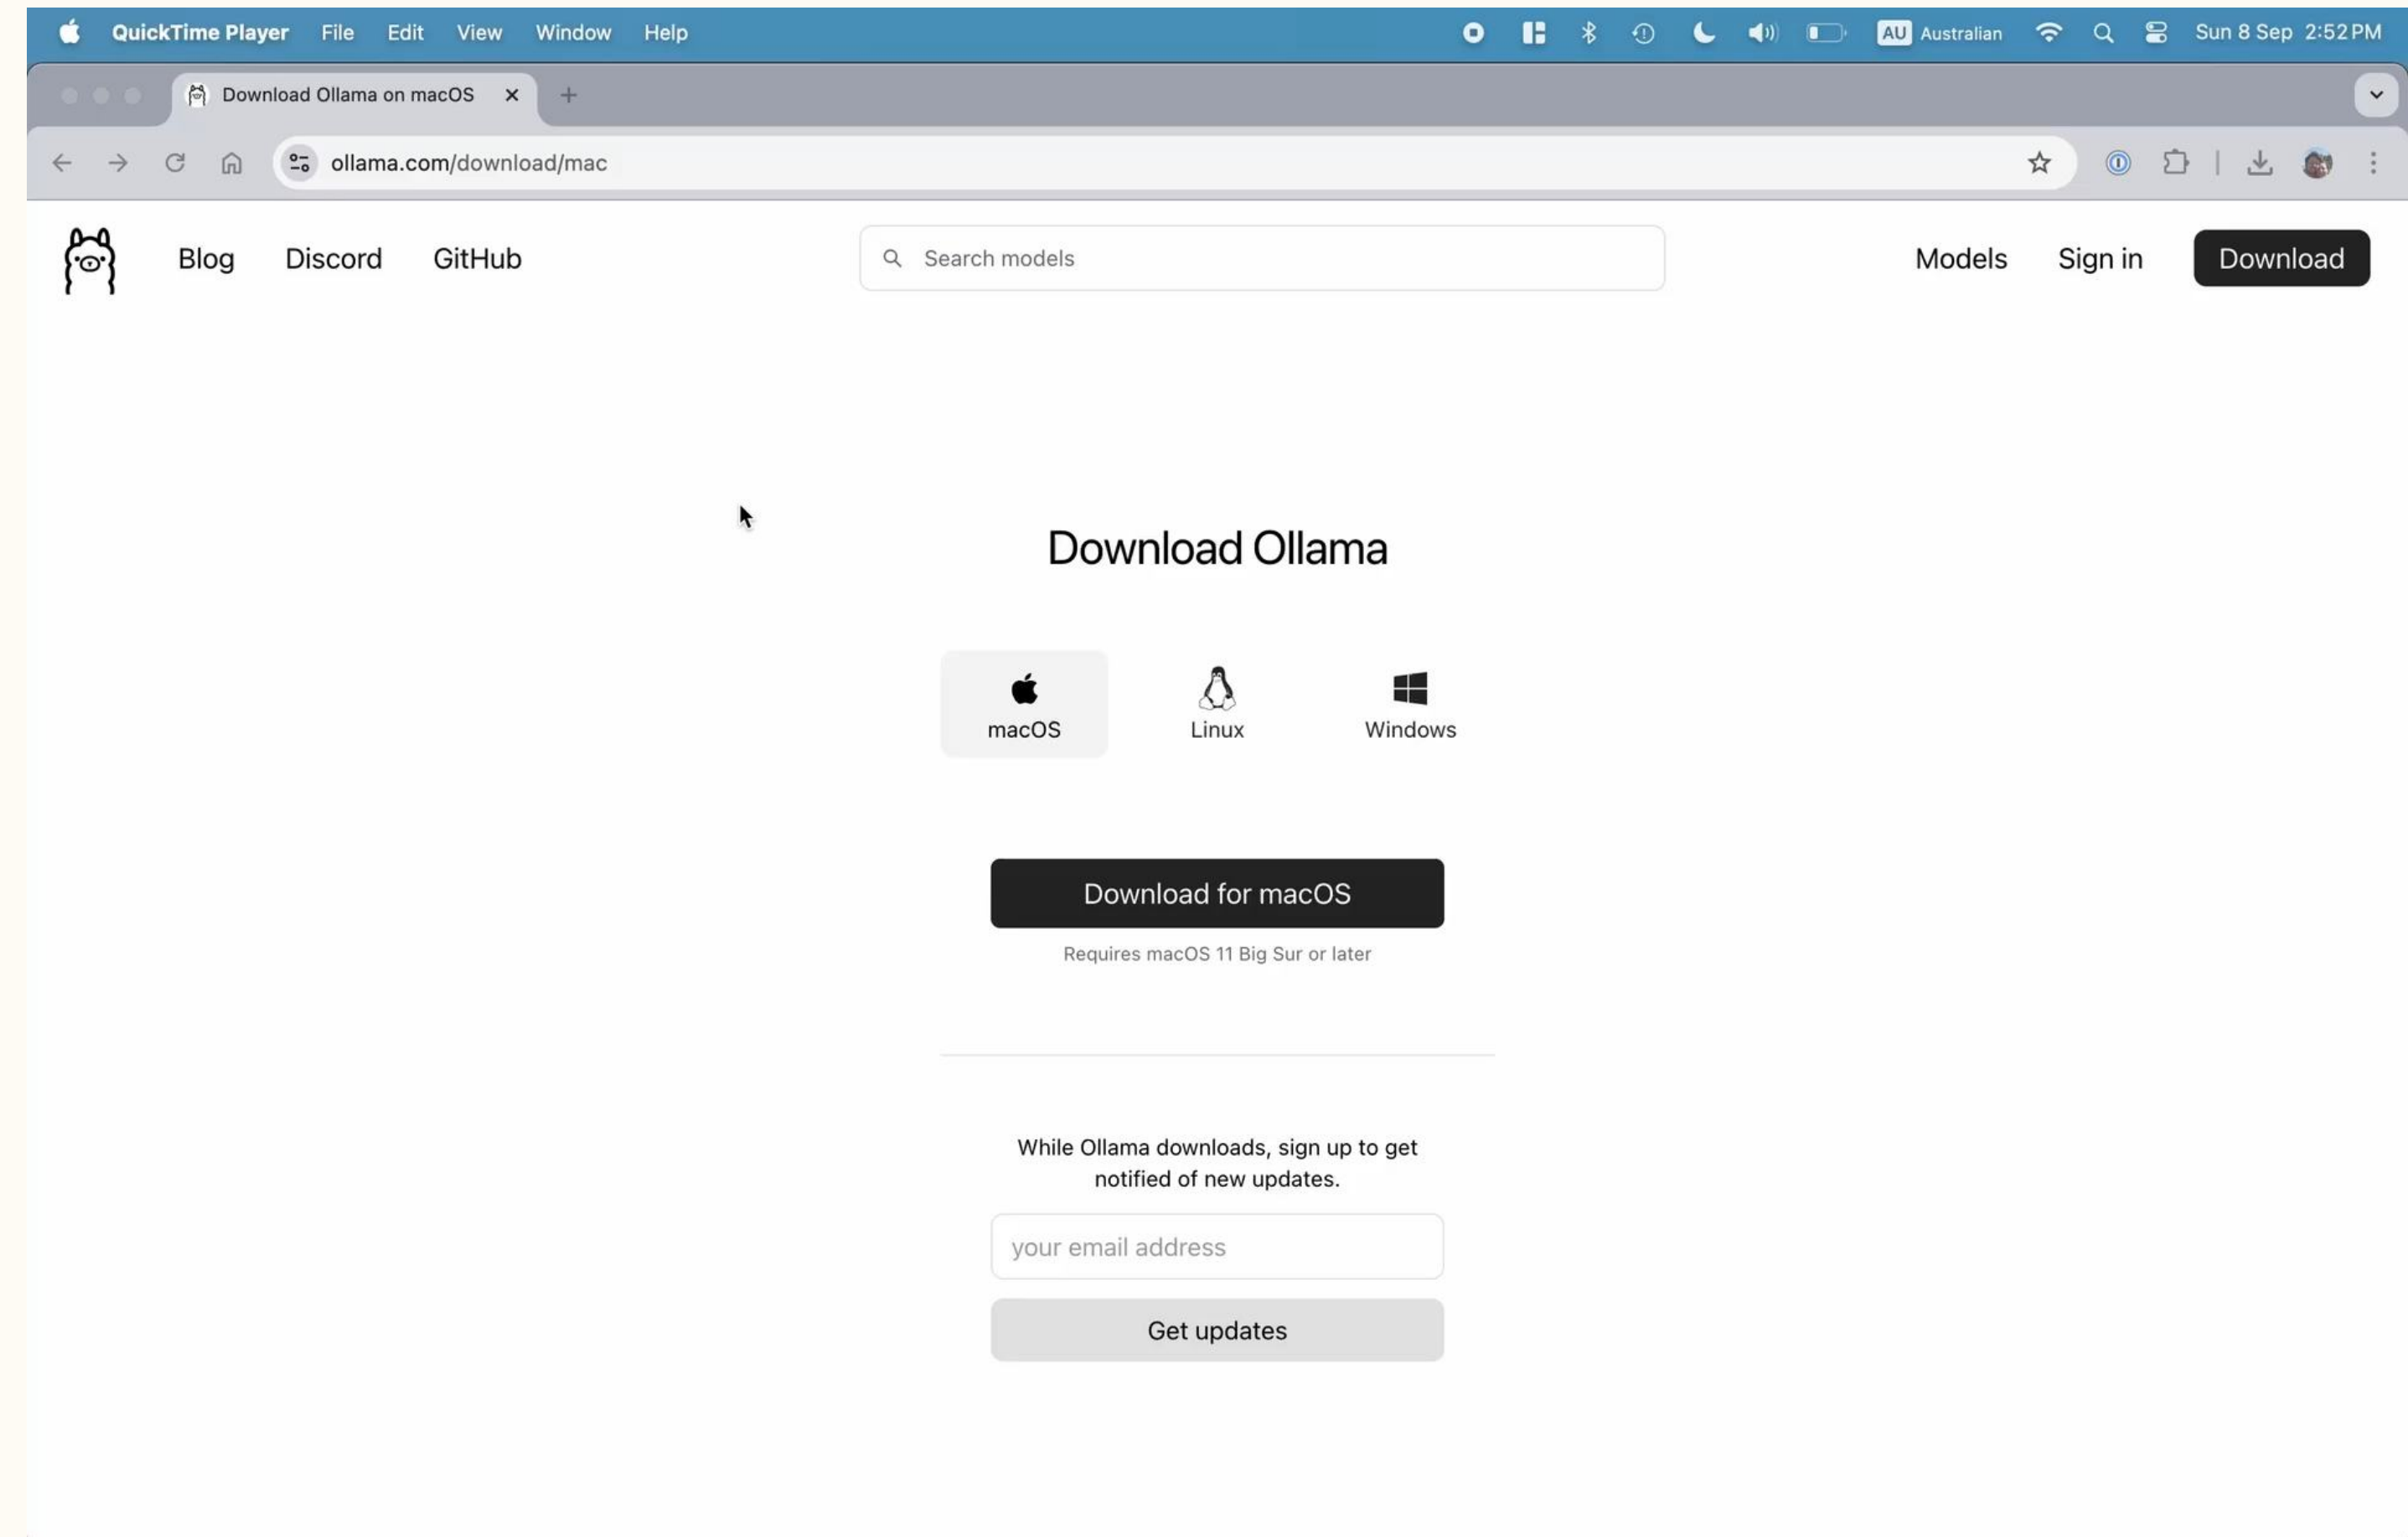
Access various models

Chat completions

Embeddings

REST API (OpenAI compatible)

Support LoRA for fine tuned models



# Ollama API

```
curl http://localhost:11434/v1/completions -d '{  
  "model": "llama3.1",  
  "prompt": "What is the meaning of life?"  
}'
```

```
curl http://localhost:11434/api/embeddings -d '{  
  "model": "llama3.1",  
  "prompt": "There is a story about eating the birds as you  
crossed south georgia island, tell me more"  
}'
```

PGVECTOR

# Postgres

Easy to install

Access pgvector

Familiar API

```
docker pull pgvector/pgvector:pg16
```

```
docker volume create pgvector-data
```

```
docker run --name pgvector-container -e  
POSTGRES_PASSWORD=<password> \  
-p 5432:5432 -v pgvector-data:/var/lib/postgresql/data \  
-d pgvector/pgvector:pg16
```

```
docker exec -it pgvector-container psql -U postgres
```

# Creating a table

```
postgres=# CREATE DATABASE localrag;  
CREATE DATABASE
```

```
postgres=# \c localrag  
You are now connected to database "localrag" as user  
"postgres".
```

```
localrag=# CREATE EXTENSION vector;  
CREATE EXTENSION
```

```
localrag=# CREATE TABLE big_items (id bigserial PRIMARY KEY,  
chunk text, embedding vector(4096));  
CREATE TABLE
```



Putting it all together: demo

# Fine tuning the Llama 3.1 model

- ✓ Improves the quality of the response
- ✓ Might be faster in production than rag
- ✓ Helps to control the output of the LLM for certain prompts
- ✓ Can be done using LoRA adapters with ollama

MLX

# MLX

Framework for training / deploying LLMs

Works on Apple Silicon

Open source from Apple

```
python3 -m venv mlx_env
```

```
. ./mlx_env/bin/activate
```

```
pip install mlx-lm
```

```
pip install "huggingface_hub[cli]"
```

```
huggingface-cli login
```

```
mlx_lm.lora --train --model meta-llama/Meta-Llama-3.1-8B --
```

```
data data --batch-size 2
```



```
Iter 800: Train loss 0.350, Learning Rate 1.000e-05, It/sec 4.490, Tokens/sec 1423.796, Trained Tokens 262104, Peak mem 17.683 GB
Iter 800: Saved adapter weights to adapters/adapters.safetensors and adapters/0000800_adapters.safetensors.
Iter 810: Train loss 0.396, Learning Rate 1.000e-05, It/sec 0.392, Tokens/sec 133.756, Trained Tokens 265518, Peak mem 17.683 GB
Iter 820: Train loss 0.342, Learning Rate 1.000e-05, It/sec 0.437, Tokens/sec 139.477, Trained Tokens 268709, Peak mem 17.683 GB
Iter 830: Train loss 0.383, Learning Rate 1.000e-05, It/sec 0.398, Tokens/sec 136.750, Trained Tokens 272142, Peak mem 17.683 GB
Iter 840: Train loss 0.364, Learning Rate 1.000e-05, It/sec 0.439, Tokens/sec 140.850, Trained Tokens 275352, Peak mem 17.683 GB
Iter 850: Train loss 0.341, Learning Rate 1.000e-05, It/sec 0.418, Tokens/sec 135.666, Trained Tokens 278600, Peak mem 17.683 GB
Iter 860: Train loss 0.347, Learning Rate 1.000e-05, It/sec 0.453, Tokens/sec 141.799, Trained Tokens 281730, Peak mem 17.683 GB
Iter 870: Train loss 0.388, Learning Rate 1.000e-05, It/sec 0.407, Tokens/sec 140.748, Trained Tokens 285186, Peak mem 17.683 GB
Iter 880: Train loss 0.368, Learning Rate 1.000e-05, It/sec 0.352, Tokens/sec 119.112, Trained Tokens 288568, Peak mem 17.683 GB
Iter 890: Train loss 0.339, Learning Rate 1.000e-05, It/sec 0.437, Tokens/sec 138.414, Trained Tokens 291737, Peak mem 17.683 GB
Iter 900: Train loss 0.358, Learning Rate 1.000e-05, It/sec 0.407, Tokens/sec 133.491, Trained Tokens 295020, Peak mem 17.683 GB
Iter 900: Saved adapter weights to adapters/adapters.safetensors and adapters/0000900_adapters.safetensors.
Iter 910: Train loss 0.348, Learning Rate 1.000e-05, It/sec 0.399, Tokens/sec 129.861, Trained Tokens 298272, Peak mem 17.683 GB
Iter 920: Train loss 0.338, Learning Rate 1.000e-05, It/sec 0.424, Tokens/sec 138.409, Trained Tokens 301533, Peak mem 17.683 GB
Iter 930: Train loss 0.353, Learning Rate 1.000e-05, It/sec 0.290, Tokens/sec 96.350, Trained Tokens 304854, Peak mem 17.683 GB
Iter 940: Train loss 0.342, Learning Rate 1.000e-05, It/sec 0.438, Tokens/sec 139.813, Trained Tokens 308049, Peak mem 17.683 GB
Iter 950: Train loss 0.345, Learning Rate 1.000e-05, It/sec 0.417, Tokens/sec 137.112, Trained Tokens 311336, Peak mem 17.683 GB
Iter 960: Train loss 0.338, Learning Rate 1.000e-05, It/sec 0.423, Tokens/sec 141.789, Trained Tokens 314688, Peak mem 17.683 GB
Iter 970: Train loss 0.340, Learning Rate 1.000e-05, It/sec 0.438, Tokens/sec 140.452, Trained Tokens 317896, Peak mem 17.683 GB
Iter 980: Train loss 0.338, Learning Rate 1.000e-05, It/sec 0.403, Tokens/sec 138.129, Trained Tokens 321322, Peak mem 17.683 GB
Iter 990: Train loss 0.335, Learning Rate 1.000e-05, It/sec 0.438, Tokens/sec 140.262, Trained Tokens 324522, Peak mem 17.683 GB
Iter 1000: Val loss 0.897, Val took 16.256s
Iter 1000: Train loss 0.337, Learning Rate 1.000e-05, It/sec 4.644, Tokens/sec 1531.171, Trained Tokens 327819, Peak mem 17.683 GB
Iter 1000: Saved adapter weights to adapters/adapters.safetensors and adapters/0001000_adapters.safetensors.
Saved final adapter weights to adapters/adapters.safetensors.
```



## Run with ollama


```
vi Modelfile
```

```
FROM llama3.1  
ADAPTER ./adapters  
~  
~  
:wq
```

```
ollama create commentcheck -f Modelfile  
ollama run commentcheck
```

# Thank you

robin@mdigital.co.nz



**Mdigital**

